

Siddhant Ray

Chicago, IL, USA
☎ +1-(773)-457-4156
✉ siddhant.r98@gmail.com
🌐 siddhant-ray.github.io/
<https://github.com/Siddhant-Ray> (GitHub)

Education

- 2023 – **The University of Chicago**, *PhD in Computer Science*
Present Advisor - Junchen Jiang and Nick Feamster
- 2020 – 2022 **ETH Zürich**, *MSc in Electrical Engineering and Information Technology*
Advisor - Laurent Vanbever
- 2016 – 2020 **VIT Vellore**, *B.Tech in Electronics and Communication Engineering*

Experience

- Jun 2026 – **Research Intern**, *Microsoft Research*
Sep 2026 ○ Building efficient models and inference systems for compound and coordinated physical AI workloads.
- Sep 2023 – **Graduate Research Assistant**, *Computer Science Department, The University of Chicago*
Present ○ Quality-cost optimization for LLM Agentic Systems with unified agent memory and KV Cache management.
○ Joint quality-latency optimization for Retrieval Augmented Generation (RAG) LLM systems with query-level configuration selection and memory management.
○ Transformer based models for sharp latency change prediction to enable packet-level queue management and tail-latency reduction for latency-sensitive applications.
- Jun 2025 – **Research Intern**, *Microsoft Research*
Sep 2025 ○ Joint internship with Outlook for scalable and cost-efficient LM architectures for email importance labeling.
○ Built a hybrid architecture with SLM cascades and embedding classifiers to reduce labeling COGS upto 150X, with negligible quality loss. Deployed in Outlook's production email client.
- Sep 2022 – **Cloud Networks Researcher**, *Advanced Network Architectures Lab, UPC Barcelona*
Mar 2023 ○ Developed an approximation for a Mixed-Integer Optimal Matching Algorithm for edge resource allocation to reduce execution time by 2.5-3X.
- Oct 2021 – **Graduate Research Assistant**, *Law, Economics, and Data Science Group, ETH Zurich*
Sep 2022 ○ Improved semantic labeling pipelines for sentence simplification, paraphrase mining, and topic clustering.
○ Developed NLP models to automatically analyse discourse in policy documents for actor-narrative clustering.
- May 2019 – **Software Development Intern**, *Capgemini Engineering*
July 2019 ○ Developed a K-Shortest Path Searching algorithm in an ONOS based Software Defined Layer 2 VPN.
○ Algorithm applied dynamic constraints of network resources (e.g.required edges) for path calculation.
- May 2018 – **Software Development Intern**, *BlueStacks*
July 2018 ○ Developed an ML algorithm to customize the Bluestack's display screen using past users' experiences.
○ Built an automation tool to create SVG cards and an address verification tool for Bluestack's game engine.

Publications

- 2026 **Siddhant Ray**, Xi Jiang, Jack Luo, Nick Feamster, and Junchen Jiang. *SwiftQueue: Optimizing Low-Latency Applications with Swift Packet Queuing*. In *New Ideas in Networked Systems (NINeS'26)*.
- 2026 **Siddhant Ray**, Ganesh Ananthanarayanan, Kevin Chian, Yan Guo, Cristina St Hill, Jack W. Stokes, Victor Wang, and Junchen Jiang. *Argo: Efficient Importance Labeling for Enterprise Email Systems*. *ArXiv preprint (version under submission)*.
- 2025 Jiayi Yao, Hanchen Li, Yuhan Liu, **Siddhant Ray**, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. *CacheBlend: Fast Large Language Model Serving for RAG with Cached Knowledge Fusion*. In *ACM European Conference on Computer Systems (EuroSys'25)*.
- 2025 **Siddhant Ray**, Rui Pan, Zhuohan Gu, Kuntai Du, Shaoting Feng, Ganesh Ananthanarayanan, Ravi Netravali, and Junchen Jiang. *METIS: Fast Quality-Aware RAG Systems with Configuration Adaptation*. In *ACM SIGOPS 31st Symposium on Operating Systems Principles (SOSP'25)*.
- 2024 Yuhan Liu, Hanchen Li, Yihua Cheng, **Siddhant Ray**, Yuyang Huang, Qizheng Zhang, Kuntai Du, Jiayi Yao, Shan Lu, Ganesh Ananthanarayanan, Michael Maire, Henry Hoffmann, Ari Holtzman, and Junchen Jiang. *CacheGen: KV Cache Compression and Streaming for Fast Large Language Model Serving*. In *ACM Special Interest Group on Data Communication (SIGCOMM'2024)*.

- 2024 Hanchen Li, Yuhan Liu, Yihua Cheng, **Siddhant Ray**, Kuntai Du, and Junchen Jiang. *Eloquent: A More Robust Transmission Scheme for LLM Token Streaming*. In *SIGCOMM Workshop on Networks for AI Computing (NAIC'24)*.
- 2022 Alexander Dietmüller, **Siddhant Ray**, Romain Jacob, and Laurent Vanbever. *A New Hope for Network Model Generalization*. In *21st ACM Workshop on Hot Topics in Networks (Hotnets'22)*.
- 2020 **Siddhant Ray** and Budhaditya Bhattacharyya. *Machine Learning based Cell Association for mMTC 5G Communication Networks*. *International Journal of Mobile Network Design and Innovation*.

Honors and Awards

- 2026 **I-Corps Cohort**, *Polisky Center and Booth School*, University of Chicago
- 2025 **Travel Grant**, *LDOS PhD Research School*, UT Austin
- 2025 **Best Paper Award**, *EuroSys 2025*, Rotterdam
- 2023 – 2028 **Liew Family Graduate Fellowship**, University of Chicago
- 2022 **Winner at Datathon**, *Microsoft Challenge*, ETH Zurich
- 2020 **Best Outgoing Student**, *SENSE department*, VIT Vellore
- 2016 – 2019 **Merit Scholarship for Academic Excellence**, VIT Vellore

Patents

- Mar 2026 **Document Labeling With Language Model Cascade**, *Microsoft*, Pending

Talks

- April 2026 **Seminar Talk**, *NAISS Lab Seminar*, HKU
- Nov 2025 **Seminar Talk**, *Systems Research Seminar*, UIUC
- Nov 2025 **Guest Lecture**, *Systems for LLMs and AI Agents Class*, UCSD
- Aug 2025 **PhD Research Talk**, *Networking Research Group*, Microsoft Research
- April 2024 **Lightning Talk**, *Research Social*, Conviva

Service

- 2026 **External Review Committee**, *MLSys*
- 2026 **Reviewer**, *AAAI, ICLR, NeurIPS*
- 2025 **Reviewer**, *ICML*
- 2025 **Artifact Evaluation Committee**, *ATC, OSDI, CoNEXT*
- 2024 **Artifact Evaluation Committee**, *CoNEXT*

Teaching

- 2026 **Teaching Assistant**, *Systems Programming II*, UChicago
- 2025 **Teaching Assistant**, *Systems Programming II*, UChicago

Other Research Projects

- 2022 Advancing Packet-Level Traffic Predictions with Transformers (Master Thesis) - [code, thesis]
- 2021 Towards a New Framework for Integration of Network Planes (Research Project) - [code]
- 2021 Attentive Neural Networks for News Classification (Research Project) - [code]

Technical Skills

- Programming Python, C++, C, Bash, Rust, SQL, Java, T_EX
- Software Linux, Git, Docker, P4, Azure, Google Cloud, AWS, Maven, K8s, Helm
- Frameworks vLLM, Langchain, Llama-Index, PyTorch, Sklearn, NLTK, Flask, Asyncio, NS3, Mininet, FRRouting