

Siddhant Ray

Chicago, IL, USA
+1-(773)-457-4156
siddhant.r98@gmail.com
<https://siddhant-ray.github.io/>
<https://github.com/Siddhant-Ray> (GitHub)

Education

- 2023 – 2028 **The University of Chicago**, *PhD in Computer Science*
Advisor - Junchen Jiang and Nick Feamster
- 2020 – 2022 **ETH Zürich**, *MSc in Electrical Engineering and Information Technology*
Advisor - Laurent Vanbever
- 2016 – 2020 **VIT Vellore**, *B.Tech in Electronics and Communication Engineering*

Experience

- Jun 2025 – **Research Intern**, *Microsoft Research*
Sep 2025
 - Joint internship with Outlook for scalable and cost-efficient LM architectures for email importance labeling.
 - Developed a hybrid architecture using SLM cascades and embedding based classifiers to reduce COGS for labeling upto 150X, with negligible quality loss.
- Sep 2023 – **Graduate Research Assistant**, *Computer Science Department, The University of Chicago*
Present
 - Project 2*: Joint quality-latency optimization for Retrieval Augmented Generation(RAG) LLM systems with query-level configuration selection and scheduling.
 - Project 1*: Developing Transformer based models for sharp latency change prediction to enable packet-level queue management for tail-latency reduction.
- Sep 2022 – **Cloud Networks Researcher**, *Advanced Network Architectures Lab, UPC Barcelona*
Mar 2023
 - Developed an approximation for a Mixed-Integer Optimal Matching Algorithm for edge resource allocation to reduce execution time by 2.5-3X.
- Oct 2021 – **Graduate Research Assistant**, *Law, Economics, and Data Science Group, ETH Zurich*
Sep 2022
 - Project 2*: Improved semantic labelling for text corpora using new NLP models, sentence simplification, paraphrase mining and clustering for topic modelling.
 - Project 1*: Creating NLP (RoBERTa based) models to analyse political discourse in meat policy documents to enable actor-narrative clustering.
- May 2019 – **Software Development Intern**, *Capgemini Engineering*
July 2019
 - Developed a K-Shortest Path Searching algorithm in an ONOS based Software Defined Layer 2 VPN.
 - Algorithm applied dynamic constraints of network resources (e.g.required edges) for path calculation.
- May 2018 – **Software Development Intern**, *BlueStacks*
July 2018
 - Developed a ML algorithm to customize the Bluestack's display screen using past users' experiences.
 - Built a automation tool for generating SVG cards for Bluestack's game engine and an address verification tool with the EasyPost API.

Publications

- 2025 Jiayi Yao, Hanchen Li, Yuhan Liu, **Siddhant Ray**, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. *CacheBlend: Fast Large Language Model Serving for RAG with Cached Knowledge Fusion*. In *Proceedings of the ACM EuroSys 2025 Conference*.
- 2025 **Siddhant Ray**, Rui Pan, Zhuohan Gu, Kuntai Du, Shaoting Feng, Ganesh Ananthanarayanan, Ravi Netravali, and Junchen Jiang. *METIS: Fast Quality-Aware RAG Systems with Configuration Adaptation*. In *Proceedings of the ACM SIGOPS 31st Symposium on Operating Systems Principles*.
- 2024 **Siddhant Ray**, Xi Jiang, Jack Luo, Nick Feamster, and Junchen Jiang. *SwiftQueue: Optimizing Low-Latency Applications with Swift Packet Queuing*, 2024. In Submission.

- 2024 Yuhan Liu, Hanchen Li, Yihua Cheng, **Siddhant Ray**, Yuyang Huang, Qizheng Zhang, Kuntai Du, Jiayi Yao, Shan Lu, Ganesh Ananthanarayanan, Michael Maire, Henry Hoffmann, Ari Holtzman, and Junchen Jiang. *CacheGen: KV Cache Compression and Streaming for Fast Large Language Model Serving*. In *Proceedings of the ACM SIGCOMM 2024 Conference*.
- 2024 Hanchen Li, Yuhan Liu, Yihua Cheng, **Siddhant Ray**, Kuntai Du, and Junchen Jiang. *Eloquent: A More Robust Transmission Scheme for LLM Token Streaming*. In *Proceedings of the 2024 SIGCOMM Workshop on Networks for AI Computing*.
- 2022 Alexander Dietmüller, **Siddhant Ray**, Romain Jacob, and Laurent Vanbever. *A New Hope for Network Model Generalization*. In *Proceedings of the 21st ACM Workshop on Hot Topics in Networks*.
- 2020 **Siddhant Ray** and Budhaditya Bhattacharyya. *Machine Learning based Cell Association for mMTC 5G Communication Networks*. *International Journal of Mobile Network Design and Innovation*, 10(1):10–16.

Honors and Awards

- 2025 **Travel Grant**, LDOS PhD Research School, UT Austin
- 2025 **Best Paper Award**, EuroSys 2025, Rotterdam
- 2023 – 2028 **Liew Family Graduate Fellowship**, University of Chicago
- 2022 **Winner at Datathon**, Microsoft Challenge, ETH Zurich
- 2020 **Best Outgoing Student**, SENSE department, VIT Vellore
- 2016 – 2019 **Merit Scholarship for Academic Excellence**, VIT Vellore

Service

- 2026 **Reviewer**, AAAI, ICLR
- 2025 **Reviewer**, ICML
- 2025 **Artifact Evaluation Committee**, ATC, OSDI, CoNEXT
- 2025 **Teaching Assistant**, CS144: Systems Programming II, UChicago
- 2024 **Artifact Evaluation Committee**, CoNEXT

Other Projects

- 2022 Advancing Packet-Level Traffic Predictions with Transformers (Master Thesis) - [code, thesis]
- 2021 Towards a New Framework for Integration of Network Planes (Research Project) - [code]
- 2021 Attentive Neural Networks for News Classification (Research Project) - [code]

Technical Skills

- Programming Python, C++, C, Bash, Rust, SQL, Java, T_EX
- Software Linux, Git, Docker, P4, Azure, Google Cloud, AWS, Maven, K8s, Helm
- Frameworks vLLM, Langchain, Llama-Index, PyTorch, Sklearn, NLTK, Flask, Asyncio, NS3, Mininet, FRRouting

Relevant Courses

- Graduate Approximation Algorithms, Algorithms, Advanced Computer Networks, System Security, Network Security, Distributed Computing, Discrete Event Systems, Networks Seminar, OS Seminar, Deep Learning, Learning Theory, Mathematics of Data Science, Neural Network Theory, Complexity Theory
- Undergraduate Computer Networks, Operating Systems, Wireless Communication, Linear Algebra

Leadership and Volunteering

- 2019 – 2020 **Technical Advisor**, IETE VIT
- 2018 – 2019 **Organizer**, TEDx VIT Vellore

