

Siddhant Ray

Chicago, IL, USA
+1-(773)-457-4156
siddhant.r98@gmail.com
www.linkedin.com/in/siddhant-ray
https://github.com/Siddhant-Ray (GitHub)

Education

- 2023 – 2028 **The University of Chicago**, *PhD in Computer Science*
Advisor - Junchen Jiang and Nick Feamster
- 2020 – 2022 **ETH Zürich**, *MSc in Electrical Engineering and Information Technology*
Advisor - Laurent Vanbever
- 2016 – 2020 **VIT Vellore**, *B.Tech in Electronics and Communication Engineering*

Experience

- Sep 2023 – **Graduate Research Assistant**, *Computer Science Department, The University of Chicago*
Present
 - *Project 1*: Developing Transformer based models for sharp latency change prediction to enable packet-level queue management for tail-latency reduction.
 - *Project 2*: Joint quality-latency optimization for Retrieval Augmented Generation(RAG) LLM systems by combining query-level configuration selection and resource scheduling.
- Sep 2022 – **Cloud Networks Researcher**, *Advanced Network Architectures Lab, UPC Barcelona*
Mar 2023
 - Reinforcement learning based resource sharing, offloading and allocation for cloud-edge systems.
 - Developed an approximation for a Mixed-Integer Optimal Matching Algorithm for edge resource allocation to reduce execution time by 2.5-3X.
- Oct 2021 – **Graduate Research Assistant**, *Law, Economics, and Data Science Group, ETH Zurich*
Sep 2022
 - *Project 1*: Improving semantic labelling for text corpora using new NLP models, sentence simplification, paraphrase mining and clustering for topic modelling.
 - *Project 2*: Creating NLP (RoBERTa based) models to analyse political discourse in meat policy documents to enable actor-narrative clustering.
- May 2019 – **Software Development Intern**, *Capgemini Engineering*
July 2019
 - Developed a K-Shortest Path Searching algorithm in an ONOS based Software Defined Layer 2 VPN.
 - Algorithm applied dynamic constraints of network resources (e.g.required edges) to be used for path calculation.
- May 2018 – **Software Development Intern**, *BlueStacks*
July 2018
 - Worked on a machine learning algorithm to predict the App Engine's appropriate display screen based on the customer's past experiences.
 - Developed an automation script for generating SVG cards for the App Engine's game front end and an address verification tool using the EasyPost API.

Publications

- 2024 Jiayi Yao, Hanchen Li, Yuhua Liu, **Siddhant Ray**, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. *CacheBlend: Fast Large Language Model Serving for RAG with Cached Knowledge Fusion*, 2024. To appear at EuroSys 2025.
- 2024 **Siddhant Ray**, Xi Jiang, Jack Luo, Nick Feamster, and Junchen Jiang. *SwiftQueue: Optimizing Low-Latency Applications with Swift Packet Queuing*, 2024. In Submission.
- 2024 Yuhua Liu, Hanchen Li, Yihua Cheng, **Siddhant Ray**, Yuyang Huang, Qizheng Zhang, Kuntai Du, Jiayi Yao, Shan Lu, Ganesh Ananthanarayanan, Michael Maire, Henry Hoffmann, Ari Holtzman, and Junchen Jiang. *CacheGen: KV Cache Compression and Streaming for Fast Large Language Model Serving*. In *Proceedings of the ACM SIGCOMM 2024 Conference*, 2024.

- 2024 Hanchen Li, Yuhua Liu, Yihua Cheng, **Siddhant Ray**, Kuntai Du, and Junchen Jiang. *Eloquent: A More Robust Transmission Scheme for LLM Token Streaming*. In *Proceedings of the 2024 SIGCOMM Workshop on Networks for AI Computing*, 2024.
- 2022 Alexander Dietmüller, **Siddhant Ray**, Romain Jacob, and Laurent Vanbever. *A New Hope for Network Model Generalization*. In *Proceedings of the 21st ACM Workshop on Hot Topics in Networks*, 2022.
- 2020 **Siddhant Ray** and Budhaditya Bhattacharyya. *Machine Learning based Cell Association for mMTC 5G Communication Networks*. *International Journal of Mobile Network Design and Innovation*, 10(1):10–16, 2020.

Technical Skills

Programming Python, C++, C, Bash, Rust, SQL, Java, T_EX
Software Linux, Git, Docker, P4 switches, ONOS, Google Cloud, AWS, Maven, MATLAB, NetSim, Cadence
Frameworks Mininet, FRRouting, PyTorch, TensorFlow, Sklearn, NLTK, Flask, SciPy, Scapy, BS4, NS-3, Langchain, vLLM

Selected Projects

- 2022 Advancing Packet-Level Traffic Predictions with Transformers (Master Thesis) - [code, thesis]
- 2021 Towards a New Framework for Integration of Network Planes (Research Project) - [code]
- 2021 Attentive Neural Networks for News Classification (Research Project) - [code]
- 2021 Investigating Possible Inductive Biases in Local Sparse Attention ViT Architectures Against Traditional CNNs (Course Project) - [code, paper]
- 2020 Maximizing Cross Traffic Flows in a L2/L3 Network with Programmable Switches (Course Project) - [code, poster]
- 2020 Machine Learning based Cell Association for 5G Communication Networks (Bachelor Thesis) - [code]

Relevant Courses

Graduate Approximation Algorithms, Algorithms, Advanced Computer Networks, System Security, Network Security, Distributed Computing, Discrete Event Systems, Networks Seminar, Deep Learning, Learning Theory, Mathematics of Data Science, Neural Network Theory
Undergraduate Computer Networks, Operating Systems, Wireless Communication, Linear Algebra

Honors and Awards

- 2023 – 2028 **Liew Family Graduate Fellowship**, University of Chicago
- 2022 **Winner at Datathon**, *Microsoft Challenge*, ETH Zurich
- 2020 **Best Outgoing Student**, *SENSE department*, VIT Vellore
- 2019 **Runner-Up at VIT Hack**, *Education Track*, VIT Vellore
- 2016 – 2019 **Merit Scholarship for Academic Excellence**, VIT Vellore

Service

CoNEXT'24 **Artifacts Evaluation Committee**

Leadership and Volunteering

- 2019 – 2020 **Technical Advisor**, IETE VIT
- 2018 – 2019 **Organizer**, TEDx VIT Vellore
- 2017 – 2020 **President** (2018 – 2019) & **Outreach Worker**, Anokha NGO